

```
#!/usr/bin/env python
```

```
.....
```

```
Vector addition using PyOpenCL.
```

```
.....
```

```
import time
```

```
import pyopencl as cl
```

```
import pyopencl.array
```

```
import numpy as np
```

```
# Select the desired OpenCL platform; you shouldn't need to change this:
```

```
NAME = 'NVIDIA CUDA'
```

```
platforms = cl.get_platforms()
```

```
devs = None
```

```
for platform in platforms:
```

```
    if platform.name == NAME:
```

```
        devs = platform.get_devices()
```

```
# Set up a command queue; we need to enable profiling to time GPU operations:
```

```
ctx = cl.Context(devs)
```

```
queue = cl.CommandQueue(ctx,
```

```
properties=cl.command_queue_properties.PROFILING_ENABLE)
```

```
# Define the OpenCL kernel you wish to run; most of the interesting stuff you
```

```
# will be doing involves modifying or writing kernels:
```

```
kernel = """ #pragma OPENCL EXTENSION cl_khr_fp64 : enable
```

```
__kernel void func( int a, int b, __global double* x, __global double* c) {
```

```
    unsigned int i = get_global_id(0);
```

```
        c[i] = c[i]+a*cos(x[i])+b*sin(x[i]);
```

```
}
```

```
.....
```

```
# Load some random data to process. Note that setting the data type is
```

```
# important; if your data is stored using one type and your kernel expects a
```

```
# different type, your program might either produce the wrong results or fail to
```

```
# run. Note that Numerical Python uses names for certain types that differ from
```

```
# those used in OpenCL. For example, np.float32 corresponds to the float type in
```

```
# OpenCL:
```

```
N = 4
```

```
a_max = 5
```

```

b_max = 7
a = np.random.randint(0,a_max,N)
b = np.random.randint(0,b_max,N)
dx = 0.05
x_max = 6.0
x = np.arange(0.0,x_max,dx)
xcopy = x
k=np.zeros_like(x)
# We can use PyOpenCL's Array type to easily transfer data from numpy arrays to
# GPU memory (and vice versa):
a_gpu = cl.array.to_device(queue, a)
b_gpu = cl.array.to_device(queue, b)
k_gpu = cl.array.to_device(queue, k)
c_gpu = cl.array.zeros(queue, x.shape, x.dtype)

# Launch the kernel; notice that you must specify the global and locals to
# determine how many threads of execution are run. We can take advantage of Numpy to
# use the shape of one of the input arrays as the global size. Since our kernel
# only accesses the global work item ID, we simply set the local size to None:
prg = cl.Program(ctx, kernel).build()

for i in xrange(0,N):
    x_gpu = cl.array.to_device(queue,x)
    prg.func(queue, x.shape, None, np.uint32(a[i]), np.uint32(b[i]), x_gpu.data, c_gpu.data)
    x = x+xcopy

# Retrieve the results from the GPU:
c = c_gpu.get()

y=np.zeros_like(x)

x = xcopy
for i in xrange(0, N):
    y += a[i]*np.cos((i+1)*x)+b[i]*np.sin((i+1)*x)

print 'input (a): ', a
print 'input (b): ', b
print 'numpy (a+b): ', y
print 'openc1 (a+b): ', c

# Compare the results from the GPU with those obtained using Numerical Python;
# this should print True:
print 'equal: ', np.allclose(y, c)

```

**# Here we compare the speed of performing the vector addition with Python and
PyOpenCL. Since the execution speed of a snippet of code may vary slightly at
different times depending on what other things the computer is running, we run
the operation we wish to time several times and average the results:**

```
M = 4
times = []
for f in xrange(M):
    start = time.time()
    for i in xrange(0, N):
        y += a[i]*np.cos((i+1)*x)+b[i]*np.sin((i+1)*x)
    times.append(time.time()-start)
```

```
print 'python time: ', np.average(times)
```

```
times = []
for f in xrange(M):
    start = time.time()
    for i in xrange(0,N):
        x_gpu = cl.array.to_device(queue,x)
        prg.func(queue, x.shape, None, np.uint32(a[i]), np.uint32(b[i]), x_gpu.data, c_gpu.data)
        x = x+xcopy
    times.append(time.time()-start)
print 'opencl time: ', np.average(times)
```

**# Notice that the
that of the PyOpenCL code for very short arrays. This is because data
transfers between host memory and that of the GPU are relatively slow. Try
gradually increasing the number of elements in a and b up to 100000 and see
what happens.**

#References: Class notes, Khronos, StackOverflow.